

UNIVERSITÀ DEGLI STUDI DI MILANO
FACOLTÀ DI SCIENZE E TECNOLOGIE

DIPARTIMENTO DI INFORMATICA
“GIOVANNI DEGLI ANTONI”



Corso di Laurea in Informatica

PROOF SEARCH IN PROPOSITIONAL LINEAR LOGIC VIA BOOLEAN CONSTRAINTS SATISFACTION

Relatore: Prof. Alberto Momigliano
Correlatore: Prof. Camillo Fiorentini

Tesi di Laurea di
Martino D'Adda
Matr. 964827

ANNO ACCADEMICO 2023-2024

Contents

Index	i
1 Introduction	1
2 The focused calculus	4
2.1 Normalization	4
2.2 Focusing	5
2.3 Constraints	6
3 Implementation	23
3.1 Why Prolog	23
3.2 Formula transformations	25
3.3 Helper predicates	26
3.4 Focusing	27
3.4.1 Identity rules	28
3.4.2 Decide rules	28
3.5 Building the tree	30
4 Related work	31
4.1 APLL	31
4.2 llprover	32
5 Testing	33
5.1 Infrastructure	33
5.1.1 Prefix format	34
5.1.2 File formats	35
5.1.3 Formula generation	36
5.2 Benchmarking	37
6 Conclusion	40

A Example derivation	41
Bibliography	44

Chapter 1

Introduction

Sequent calculus is a formalism first introduced by G. Gentzen in 1934 [6, 7]. In its classical interpretation, a sequent is just an implication between finite (possibly empty) sequences of formulae, such that

$$\Delta \vdash \Gamma$$

is to be interpreted as

$$\bigwedge_{\phi \in \Delta} \phi \Rightarrow \bigvee_{\gamma \in \Gamma} \gamma$$

and Δ is called the antecedent, and Γ the succedent.

Sequents are manipulated using rules that have sequents as premises and as a conclusion. These rules are used to describe connectives operationally: defining their behavior by how they are derived from smaller formulae. A rule acting on the antecedent is called a left rule whereas a rule acting on the succedent is called a right rule. A proof is then a tree where the root is the sequent to prove, and the inner nodes are the rules chained one to another by matching premises to conclusions. This way the leaves are just rules without any premises, also called axioms. In this context, a theorem prover – or simply a prover – is an algorithm that tries to build a proof of a given sequent to determine if it is true or not. More precisely a bottom-up prover starts at the final sequent, and tries to build the proof tree by applying the rules of the calculus backwards, breaking down the formulae into their subformulae.

There are three special rules – called structural rules – often left implicit in proofs. We give their left versions:

- weakening: one may “weaken” the sequent by adding a proposition without changing the truth of the sequent

$$\frac{\Gamma \vdash \Delta}{\Gamma, \phi \vdash \Delta}$$

- contraction: one may “contract” two copies of the same proposition into one without changing the truth of the sequent

$$\frac{\Gamma, \phi, \phi \vdash \Delta}{\Gamma, \phi \vdash \Delta}$$

- exchange: one may switch position of two propositions in a sequent freely without changing its truth

$$\frac{\Gamma, \phi, \psi \vdash \Delta}{\Gamma, \psi, \phi \vdash \Delta}$$

These do not directly specify the behavior of some connective, but rather show the ways in which a formula may be moved inside the sequent. Logics missing at least one of these three rules are called *substructural*; it is in this class that lies linear logic.

Linear logic is a logic proposed by J.-Y. Girard in his seminal paper of 1987 [8]. The distinctive trait of this logic is that in general its formulae cannot be copied or discarded, but instead are consumed. Put differently, its sequent calculus lacks the structural rules of weakening and contraction¹.

Under these assumptions a certain sequent (here we are interpreting sequents as multisets) is provable if and only if all of its formulae get used exactly once; for this reason this logic is sometimes called a logic of resources, in the same way classical logic is a logic of truths and intuitionistic logic is a logic of proofs.

In linear logic each connective from classical logic has two interpretations: an additive one – where the context is the same for all the premises, and a multiplicative one – where the context is partitioned between the premises. To better understand why, we analyze two rules for classic conjunction:

$$\frac{\Delta \vdash \phi_2, \Gamma \quad \Delta \vdash \phi_1, \Gamma}{\Delta \vdash \phi_1 \wedge \phi_2, \Gamma} \quad \frac{\Delta' \vdash \phi_1, \Gamma' \quad \Delta'' \vdash \phi_2, \Gamma''}{\Delta', \Delta'' \vdash \phi_1 \wedge \phi_2, \Gamma', \Gamma''}$$

The two rules above are equivalent only if the use of weakening and contraction is permitted; for this reason in linear logic the two interpretations are distinct: the left one is the additive one, and the right one the multiplicative one. The same holds for the constants \top and \perp , which also have two versions. Table 1 shows the linear logic connectives corresponding to the classical ones.

Linear logic defined as of right now, albeit having the added complexity of splitting, is nonetheless decidable: since formulae are finite and they cannot be copied, it is possible to explore all the possibilities. Undecidability is reintroduced by localizing the uses of contraction and weakening by means of two modalities – “of-course”,

¹There are logics where even the exchange rule is missing, these are called non-commutative because the order of the formulae in the sequent matters (e.g. Lambek calculus).

Class.	Add.	Mult.
\wedge	$\&$	\otimes
\vee	\oplus	\wp
\top	\top	1
\perp	0	\perp

Table 1: Classical connectives and their corresponding linear ones.

written $!\phi$ and “why-not”, written $?\phi$ – called exponentials.

It is the multiplicative side which brings the most complexity during bottom-up proof search: the action of partitioning the sequent – called splitting – may imply an exponential number of attempts to find the correct subsets. Some methods have been proposed to alleviate this, such as the input/output method [10].

In 2001 D. Pym and J. Harland publish a paper [9] where they present a new way of tackling the problem of splitting by means of boolean constraints. These constraints are generated during proof search from boolean expressions associated to the formulae, and are used to enforce linearity. This way the complexity shifts from choosing the right set of formulas to prove a certain branch, to solving for boolean assignment – a problem for which there are much more sophisticated algorithms.

This thesis is organized as follows:

- in §2 we define a focused and one-sided version to the calculus described in [9], and we give a proof of its soundness consisting of a forgetful functor to the triadic calculus of [2];
- in §3 we discuss a Prolog implementation of the calculus of the previous chapter;
- in §4 we quickly describe the main implementation details of two other bottom-up provers for full linear logic: *llprover* [15] and *APLL* [18];
- finally in §5 we present the framework built to test and compare our prover with the others from the previous section, and then we show the results of these benchmarks.

Chapter 2

The focused calculus

In this chapter we will define a focused one sided constraint calculus for full linear logic. This calculus is a hybrid between the one defined in [9] and the triadic calculus of [2]. Before all, we give the definition of a linear logic formula.

Definition 2.0.1 (Linear logic formula). A propositional linear logic formula is defined as follows:

$$\begin{array}{rcl}
 \phi ::= & \phi \otimes \phi & | \quad 1 \\
 & \phi \wp \phi & | \quad \perp \\
 & \phi \multimap \phi & \\
 & \phi \oplus \phi & | \quad 0 \\
 & \phi \& \phi & | \quad \perp \\
 & !\phi & | \quad ?\phi \\
 & \phi^\perp & \\
 & \alpha &
 \end{array}
 \left. \begin{array}{l} \\ \\ \\ \\ \\ \\ \end{array} \right\} \begin{array}{l} \text{(Multiplicatives and their identities)} \\ \\ \text{(Additives and their identities)} \\ \text{(Exponentials)} \\ \text{(Negation)} \\ \text{(Atom)} \end{array}$$

We will use ϕ to denote formulae, α to denote atoms, and Δ or Γ to denote multisets of formulae.

2.1 Normalization

Since in linear logic negation is symmetric and an involution, it is usual to work only with formulae in negated normal form.

Definition 2.1.1 (Negated Normal Form – NNF). A formula is in NNF if all its linear implications (\multimap) are expanded to pars (\wp) using the tautology

$$a \multimap b \Leftrightarrow a^\perp \wp b$$

and negation is pushed down to atoms. Alternatively a formula in NNF is defined as follows:

$$\phi ::= \begin{array}{l|l} \phi \otimes \phi & 1 \\ \phi \wp \phi & \perp \\ \phi \oplus \phi & 0 \\ \phi \& \phi & \perp \\ !\phi & ?\phi \\ \alpha^\perp & \alpha \end{array}$$

Accordingly, a sequent is in NNF iff all of its formulae are in NNF.

A generic formula is normalized by applying recursively the DeMorgan rules for linear logic until NNF is reached. Normalization of judgments instead takes a two-sided judgment of the form

$$\Delta \vdash \Gamma$$

and transforms it into a one-sided judgment

$$\vdash \Delta'$$

where the right side is composed of the normalization of Γ and Δ^\perp .

Normalization has some implementation-wise advantages, but for now it is only important because it shrinks the size of the complete calculus by roughly a half since only the right rules remain, thus making the calculus easier to represent (see. Figure 2 and 3) and making the proof of Theorem 2.3.1 significantly shorter.

2.2 Focusing

Focusing is a technique described by Andreoli in his seminal paper [2]. In it he recognizes two alternating phases in a proof: a deterministic phase – called asynchronous phase; and a non-deterministic phase – called synchronous phase or focused phase. In the asynchronous phase the applicable rules are all invertible, making their order of application irrelevant. For this reason these rules are called asynchronous rules, whereas the non-invertible ones are called synchronous. In this system to each formula is assigned a positive or negative polarity based on its top-level connective:

- connectives with a synchronous right rule are defined to have a positive polarity, these are:

$$\otimes, \oplus, !, 1$$

- whereas connectives with an asynchronous right rule have a negative polarity, these are:

$$\wp, \&, ?, \top, \perp$$

For atoms polarities may be assigned with some arbitrarily complex mechanisms. We will follow [11] and simply assign atoms with a negative polarity and negated atoms with a positive one. Since we work in one sided linear logic and connectives have only right rules, positive connectives may also be called synchronous, and negative connectives asynchronous.

Definition 2.2.1. Based on the definitions above we get the following predicates:

- “ ϕ atom” is true whenever ϕ is an atom;
- “ ϕ asy” is true whenever ϕ has as an asynchronous top-level connective;
- “ ϕ negative” is true whenever ϕ is either an atom or an asynchronous connective, so

$$\phi \text{ negative} = \phi \text{ atom} \vee \phi \text{ asy}$$

Figure 1 shows the focused calculus described by Anderoli in [2]. This calculus is also called triadic calculus, since its judgments have three members: a set of unrestricted formulae; a multiset of linear formulae put to the side; and either a single formula or another multiset of linear formulae depending on the current phase. An arrow pointing up symbolizes the asynchronous phase, and an arrow pointing down the focusing phase. Phase switching happens using the decide rules $([D_1]_A, [D_2]_A)$, which non-deterministically choose a formula to focus on.

2.3 Constraints

During proof search constraints are generated as the formulae get broken up. These are created in a particular way, as to guarantee that if an assignment exists that satisfies them, then the formulae have been used linearly.

Definition 2.3.1 (Variables, expressions). A boolean variable is simply a symbol to which one can associate a value of true or false. A boolean expression, in our case, is just a conjunction of possibly negated boolean variables as follows

$$e ::= x \wedge e \mid \bar{x} \wedge e \mid x$$

We will call e such a conjunction and x the single boolean variables. Given a boolean expression e we write

$$\text{vars}(e) = \{x_i \mid x_i \in e\}$$

as the set of the variables appearing in the expression.

$$\begin{array}{c}
[\perp]_A \frac{\vdash_A \Psi : \Delta \uparrow \Gamma}{\vdash_A \Psi : \Delta \uparrow \perp, \Gamma} \qquad [\top]_A \frac{}{\vdash_A \Psi : \Delta \uparrow \top, \Gamma} \\
[\wp]_A \frac{\vdash_A \Psi : \Delta \uparrow \phi_1, \phi_2, \Gamma}{\vdash_A \Psi : \Delta \uparrow \phi_1 \wp \phi_2, \Gamma} \qquad [\&]_A \frac{\vdash_A \Psi : \Delta \uparrow \phi_1, \Gamma \quad \vdash_A \Psi : \Delta \uparrow \phi_2, \Gamma}{\vdash_A \Psi : \Delta \uparrow \phi_1 \& \phi_2, \Gamma} \\
[\oplus_L]_A \frac{\vdash_A \Psi : \Delta \downarrow \phi_1}{\vdash_A \Psi : \Delta \downarrow \phi_1 \oplus \phi_2} \qquad [\oplus_R]_A \frac{\vdash_A \Psi : \Delta \downarrow \phi_2}{\vdash_A \Psi : \Delta \downarrow \phi_1 \oplus \phi_2} \\
[\otimes]_A \frac{\vdash_A \Psi : \Gamma \downarrow \phi_1 \quad \vdash_A \Psi : \Delta \downarrow \phi_2}{\vdash_A \Psi : \Gamma, \Delta \downarrow \phi_1 \otimes \phi_2} \qquad [1]_A \frac{}{\vdash_A \Psi : \cdot \downarrow 1} \\
[?]_A \frac{\vdash_A \phi, \Psi : \Delta \uparrow \Gamma}{\vdash_A \Psi : \Delta \uparrow ?\phi, \Gamma} \qquad [!]_A \frac{\vdash_A \Psi : \cdot \uparrow \phi}{\vdash_A \Psi : \cdot \downarrow !\phi} \\
[I_1]_A \frac{\alpha \text{ atom}}{\vdash_A \Psi : \alpha \downarrow \alpha^\perp} \qquad [I_2]_A \frac{\alpha \text{ atom}}{\vdash_A \alpha, \Psi : \cdot \downarrow \alpha^\perp} \\
[D_1]_A \frac{\phi \text{ not atom} \quad \vdash_A \Psi : \Delta \downarrow \phi}{\vdash_A \Psi : \phi, \Delta \uparrow \cdot} \qquad [D_2]_A \frac{\phi \text{ not atom} \quad \vdash_A \Psi : \Delta \downarrow \phi}{\vdash_A \phi, \Psi : \Delta \uparrow \cdot} \\
[R\uparrow]_A \frac{\phi \text{ not asy} \quad \vdash_A \Psi : \phi, \Delta \uparrow \Gamma}{\vdash_A \Psi : \Delta \uparrow \phi, \Gamma} \qquad [R\downarrow]_A \frac{\phi \text{ negative} \quad \vdash_A \Psi : \Delta \uparrow \phi}{\vdash_A \Psi : \Delta \downarrow \phi}
\end{array}$$

Figure 1: J.-M. Andreoli's triadic calculus.

Definition 2.3.2 (Annotated formula). Given a formula ϕ defined as in Definition 2.1.1 and a boolean expression e defined as in Definition 2.3.1, an *annotated formula* is simply a pair

$$\langle \phi, e \rangle$$

that associates the formula to the expression. We denote

- the operation of extracting the boolean expression associated to a given formula as

$$\text{expr}(\langle \phi, e \rangle) = e$$

and then extend this notation to sequents such that $\text{expr}(\Delta)$ is the set of all

boolean expressions of Δ

$$\text{expr}(\Delta) = \{\text{expr}(\phi) \mid \phi \in \Delta\}$$

- the operation of extracting the set of variables appearing in the expression e as

$$\text{vars}(\langle \phi, e \rangle) = \text{vars}(e)$$

and then extend this notation to sequents such that $\text{vars}(\Delta)$ is the set of all the variables appearing in the boolean expressions of the annotated formulae of Δ

$$\text{vars}(\Delta) = \{\text{vars}(\text{expr}(\phi)) \mid \phi \in \Delta\}$$

It is important to note that only the topmost connective gets annotated, and not the sub-formulae.

The purpose of putting formulae and expressions together in the annotated formula is twofold:

- the actions taken on the formula determine the constraints that will be generated, and these refer to the expressions associated to said formula;
- after the constraints are solved we can query the assignment of the variables and find out if the associated formula is used or not in a certain branch of a proof (see Definition 2.3.5).

The above constraints may be only of two kinds: “ e avail” and “ e used”.

Definition 2.3.3 (Constraints). Given an annotated formula $\langle \phi, e \rangle$ as in Definition 2.3.2, a constraint λ is either

- “ e used”, which states that the formula ϕ gets consumed in this branch of the proof;
- “ e avail”, which states that the formula ϕ does not get consumed in this branch of the proof, and thus is available to be used in another branch.

We then extend these constraints to sequents, such that

$$\begin{aligned} \Delta \text{ used} &= \{e \text{ used} \mid e \in \text{exp}(\Delta)\} \\ \Delta \text{ avail} &= \{e \text{ avail} \mid e \in \text{exp}(\Delta)\} \end{aligned}$$

We denote

- $\text{exp}(\lambda) = e$ where λ is either “ e used” or “ e avail”;

- $\text{vars}(\lambda) = \text{vars}(\text{exp}(\lambda))$, and extend this notation to a set of constraints Λ

$$\text{vars}(\Lambda) = \bigcup_{\lambda \in \Lambda} \text{vars}(\lambda)$$

Definition 2.3.4 (Assignment). An assignment is a function

$$V : \{\dots, x_i, x_j, \dots\} \rightarrow \{\top, \perp\}$$

that associates the members of a set of variables to either true or false. Given a set of variables X , we say that the assignment V *covers* X if no variable in X is left undefined under V , or

$$X \subseteq \text{Dom}(V)$$

Definition 2.3.5 (Evaluation). Given a boolean expression e and an assignment V , such that V covers $\text{vars}(e)$, we write

$$e[V] = e[\dots, x_i := V(x_i), x_j := V(x_j), \dots]$$

as the value of the expression e substituting its variables using assignment V . We extend this notation, such that

- given a constraint λ and an assignment V that covers $\text{vars}(\lambda)$

$$\begin{cases} e \text{ used}[V] = \top & \text{if } e[V] = \top \\ e \text{ used}[V] = \perp & \text{if } e[V] = \perp \\ e \text{ avail}[V] = \top & \text{if } e[V] = \perp \\ e \text{ avail}[V] = \perp & \text{if } e[V] = \top \end{cases}$$

- given an annotated sequent Δ and an assignment V that covers $\text{vars}(\Delta)$

$$\Delta[V] = \{\phi \mid \langle \phi, e \rangle \in \Delta, e[V] = \top\}$$

Definition 2.3.6. Given two assignments V' and V'' and a set of variables X , we say that V' and V'' coincide for X – written $V' \sim_X V''$ – when they both map all the variables in X to the same values, or

$$V' \sim_X V'' \Leftrightarrow \forall x \in X \mid V'(x) = V''(x)$$

Definition 2.3.7 (Satisfiability). Given a set of constraints Λ and an assignment function V that covers $\text{vars}(\Lambda)$, we say that V satisfies Λ iff every constraint in Λ is

true under V , or

$$\Lambda \downarrow V \Leftrightarrow \bigwedge_{\lambda \in \Lambda} \lambda[V] = \top$$

We now expand the concept of triadic judgment from [2] (Figure 1) by adding explicit constraints.

Definition 2.3.8. Given any judgment in our calculus, it can be in either two forms:

- focused or in the synchronous phase, written:

$$\vdash \Psi : \Delta \Downarrow \phi \parallel \Lambda : V$$

- in the asynchronous phase, written:

$$\vdash \Psi : \Delta \Uparrow \Gamma \parallel \Lambda : V$$

Where

- Ψ is the same as the triadic calculus of [2]: a set of unrestricted non annotated formulae, or all formulae that can be freely discarded or duplicated.
- Δ and Γ are multisets of linear (annotated) formulae, these are respectively the formulae “put to the side” and the formulae which are being “worked on” during a certain moment of the asynchronous phase;
- Λ and V are the constraints and the assignment as defined in Definition 2.3.7. By adding these members we make the flow of constraints through the proof tree explicit, leaving no ambiguity to where the constraints should be checked. This approach to constraints differs from the one in [9]. The choice of letters is mainly a mnemonic one, constraints Λ “go-up” the proof tree and solutions V “come down” from the leaves.

Definition 2.3.9 (Splitting). Given a sequent of annotated formulae Δ and a set of variables X such that $|\Delta| = |X|$, we define the operation of splitting it as a function

$$\text{split}_X(\Delta) \mapsto (\Delta_L, \Delta_R)$$

where

$$\begin{aligned} \Delta_L &= \{\langle \phi_i, x_i \wedge e_i \rangle \mid i \in \{1, \dots, n\}\} \\ \Delta_R &= \{\langle \phi_i, \bar{x}_i \wedge e_i \rangle \mid i \in \{1, \dots, n\}\} \end{aligned}$$

with n the cardinality of Δ , and ϕ_i (resp. e_i) the formula (resp. the expression) of the i -eth annotated formula in Δ using an arbitrary order. The same holds for x_i and X . With a slight abuse of notation we will write Δ_L^X and Δ_R^X to mean respectively the left projection and the right projection of the pair (Δ_L, Δ_R) .

As a small example for clarity, given the sequent

$$\begin{aligned}\Delta &= \langle a \otimes b, x_1 \rangle, \langle c^\perp, x_2 \rangle \\ X &= \{x_3, x_4\}\end{aligned}$$

this is split into

$$\begin{aligned}\Delta_L^X &= \langle a \otimes b, x_3 \wedge x_1 \rangle, \langle c^\perp, x_4 \wedge x_2 \rangle \\ \Delta_R^X &= \langle a \otimes b, \overline{x_3} \wedge x_1 \rangle, \langle c^\perp, \overline{x_4} \wedge x_2 \rangle\end{aligned}$$

Figure 2 shows the calculus of [9] using our notation of explicit constraints, Figure 3 shows our focused calculus instead. Pym and Harland in [9] describe a number different ways of accumulating and checking constraints, roughly comparable to breadth first search, depth first search and a third intermediate method. Figure 2 and 3 represent what [9] calls a “lazy” strategy, corresponding to a DFS. In both the calculi it is easy to see that most of the rules simply make sure the constraints remain consistent, and only the axioms ($[I_1]$, $[I_2]$, $[A]_{\text{PH}}$, ...) and the tensor ($[\otimes]$, $[\otimes]_{\text{PH}}$) actually introduce new variables or significant constraints.

Definition 2.3.10 (New variables). In Figure 2 and 3 we write

$$x \text{ new}$$

to state that the variable name x has not yet occurred in any expression of the proof tree. One can think of a counter threaded through the proof tree, which is incremented for each new variable:

$$[D_2] \frac{\phi \text{ not atom} \quad x_i \text{ new} \quad \vdash^{i+1,n} \Psi : \Delta \uparrow \langle \phi, x_i \rangle \parallel x_i \text{ used}, \Lambda : V}{\vdash^{i,n} \phi, \Psi : \Delta \uparrow . \parallel \Lambda : V}$$

here i is the current value of the counter, and n is the value of the counter exiting the proof, such that:

$$[I_1] \frac{x \text{ atom} \quad e_1 \text{ used}, e_2 \text{ used}, \Delta \text{ avail} \downarrow V}{\vdash^{i,i} \Psi : \langle \alpha, e_1 \rangle, \Delta \uparrow \langle \alpha^\perp, e_2 \rangle \parallel \Lambda : V}$$

Given this, we then define

$$X \text{ new}$$

$$\begin{array}{c}
[\mathcal{V}]_{\text{PH}} \frac{\vdash_{\text{PH}} \Delta, \langle \phi_1, e \rangle, \langle \phi_2, e \rangle \parallel e \text{ used}, \Lambda : V}{\vdash_{\text{PH}} \Delta, \langle \phi_1 \mathcal{V} \phi_2, e \rangle \parallel \Lambda : V} \\
\\
[\perp]_{\text{PH}} \frac{\vdash_{\text{PH}} \Delta \parallel e \text{ used}, \Lambda : V}{\vdash_{\text{PH}} \Delta, \langle \perp, e \rangle \parallel \Lambda : V} \qquad [\top]_{\text{PH}} \frac{}{\vdash_{\text{PH}} \Delta, \langle \top, - \rangle \parallel \Lambda : V} \\
\\
[\&]_{\text{PH}} \frac{\vdash_{\text{PH}} \Delta, \langle \phi_1, e \rangle \parallel e \text{ used}, \Lambda : V'' \quad \vdash_{\text{PH}} \Delta, \langle \phi_2, e \rangle \parallel e \text{ used}, \Lambda : V' \quad V' \sim_{\text{vars}(\Delta, \Gamma)} V''}{\vdash_{\text{PH}} \Delta, \langle \phi_1 \& \phi_2, e \rangle \parallel \Lambda : V''} \\
\\
[\otimes]_{\text{PH}} \frac{X \text{ new} \quad \vdash_{\text{PH}} \Delta_L^X, \langle \phi_1, e \rangle \parallel e \text{ used}, \Lambda : V' \quad \vdash_{\text{PH}} \Delta_R^X, \langle \phi_2, e \rangle \parallel e \text{ used}, \Lambda : V'' \quad V' \sim_{\text{vars}(\Delta), X} V''}{\vdash_{\text{PH}} \Delta, \langle \phi_1 \otimes \phi_2, e \rangle \parallel \Lambda : V''} \\
\\
[\oplus]_{\text{PH}} \frac{x \text{ new} \quad \vdash_{\text{PH}} \Delta, \langle \phi_1, x \rangle, \langle \phi_2, \bar{x} \rangle \parallel e \text{ used}, \Lambda : V}{\vdash_{\text{PH}} \Delta, \langle \phi_1 \oplus \phi_2, e \rangle \parallel \Lambda : V} \\
\\
[1]_{\text{PH}} \frac{e \text{ used}, \Lambda, \Delta \text{ avail} \downarrow V}{\vdash_{\text{PH}} \Delta, \langle 1, e \rangle \parallel \Lambda : V} \qquad [!]_{\text{PH}} \frac{\vdash_{\text{PH}} ?\Delta, \langle \phi, e \rangle \parallel e \text{ used}, \Lambda : V}{\vdash_{\text{PH}} ?\Delta, \langle !\phi, e \rangle \parallel \Lambda : V} \\
\\
[?]_{\text{PH}} \frac{\vdash_{\text{PH}} \Delta, \langle \phi, e \rangle \parallel e \text{ used}, \Lambda : V}{\vdash_{\text{PH}} \Delta, \langle ?\phi, e \rangle \parallel \Lambda : V} \\
\\
[W?]_{\text{PH}} \frac{\vdash_{\text{PH}} \Delta \parallel e \text{ used}, \Lambda : V}{\vdash_{\text{PH}} \langle ?\phi, e \rangle, \Delta \parallel \Lambda : V} \qquad [C?]_{\text{PH}} \frac{\vdash_{\text{PH}} \langle ?\phi, e \rangle, \langle ?\phi, e \rangle, \Delta \parallel e \text{ used}, \Lambda : V}{\vdash_{\text{PH}} \langle ?\phi, e \rangle, \Delta \parallel \Lambda : V} \\
\\
[A]_{\text{PH}} \frac{\alpha \text{ atom} \quad \Lambda, e_1 \text{ used}, e_2 \text{ used}, \Delta \text{ avail} \downarrow V}{\vdash_{\text{PH}} \langle \alpha, e_1 \rangle, \langle \alpha^\perp, e_2 \rangle, \Delta \parallel \Lambda : V}
\end{array}$$

Figure 2: The one sided version of the calculus from [9] with explicit constraint propagation.

$$\begin{array}{c}
[\wp] \frac{\vdash \Psi : \Delta \uparrow \langle \phi_1, e \rangle, \langle \phi_2, e \rangle, \Gamma \parallel e \text{ used}, \Lambda : V}{\vdash \Psi : \Delta \uparrow \langle \phi_1 \wp \phi_2, e \rangle, \Gamma \parallel \Lambda : V} \\
[\perp] \frac{\vdash \Psi : \Delta \uparrow \Gamma \parallel e \text{ used}, \Lambda : V}{\vdash \Psi : \Delta \uparrow \langle \perp, e \rangle, \Gamma \parallel \Lambda : V} \quad [\top] \frac{}{\vdash \Psi : \Delta \uparrow \langle \top, - \rangle, \Gamma \parallel \Lambda : V} \\
[\&] \frac{\vdash \Psi : \Delta \uparrow \langle \phi_1, e \rangle, \Gamma \parallel e \text{ used}, \Lambda : V'' \quad \vdash \Psi : \Delta \uparrow \langle \phi_2, e \rangle, \Gamma \parallel e \text{ used}, \Lambda : V', \quad V' \sim_{\text{vars}(\Delta, \Gamma)} V''}{\vdash \Psi : \Delta \uparrow \langle \phi_1 \& \phi_2, e \rangle, \Gamma \parallel \Lambda : V''} \\
[?] \frac{\vdash \phi, \Psi : \Delta \uparrow \Gamma \parallel \Lambda : V}{\vdash \Psi : \Delta \uparrow \langle ?\phi, - \rangle, \Gamma \parallel \Lambda : V} \\
[R\uparrow] \frac{\phi \text{ not asy} \quad \vdash \Psi : \langle \phi, e \rangle, \Delta \uparrow \Gamma \parallel \Lambda : V}{\vdash \Psi : \Delta \uparrow \langle \phi, e \rangle, \Gamma \parallel \Lambda : V}
\end{array}$$

(a) Asynchronous rules.

$$\begin{array}{c}
[\otimes] \frac{X \text{ new} \quad \vdash \Psi : \Delta_L^X \Downarrow \langle \phi_1, e \rangle \parallel e \text{ used}, \Lambda : V', \quad \vdash \Psi : \Delta_R^X \Downarrow \langle \phi_2, e \rangle \parallel e \text{ used}, \Lambda : V'' \quad V' \sim_{\text{vars}(\Delta), X} V''}{\vdash \Psi : \Delta \Downarrow \langle \phi_1 \otimes \phi_2, e \rangle \parallel \Lambda : V''} \\
[\oplus_L] \frac{\vdash \Psi : \Delta \Downarrow \langle \phi_1, e \rangle \parallel e \text{ used}, \Lambda : V}{\vdash \Psi : \Delta \Downarrow \langle \phi_1 \oplus \phi_2, e \rangle \parallel \Lambda : V} \quad [\oplus_R] \frac{\vdash \Psi : \Delta \Downarrow \langle \phi_2, e \rangle \parallel e \text{ used}, \Lambda : V}{\vdash \Psi : \Delta \Downarrow \langle \phi_1 \oplus \phi_2, e \rangle \parallel \Lambda : V} \\
[1] \frac{e \text{ used}, \Delta \text{ avail}, \Lambda \Downarrow V}{\vdash \Psi : \Delta \Downarrow \langle 1, e \rangle \parallel \Lambda : V} \quad [!] \frac{\vdash \Psi : \Delta \Downarrow \langle \phi, e \rangle \parallel e \text{ used}, \Delta \text{ avail}, \Lambda : V}{\vdash \Psi : \Delta \Downarrow \langle !\phi, e \rangle \parallel \Lambda : V} \\
[R\Downarrow] \frac{\phi \text{ negative} \quad \vdash \Psi : \Delta \uparrow \langle \phi, e \rangle \parallel e \text{ used}, \Lambda : V}{\vdash \Psi : \Delta \Downarrow \langle \phi, e \rangle \parallel \Lambda : V}
\end{array}$$

(b) Synchronous rules.

$$\begin{aligned}
[I_1] & \frac{\alpha \text{ atom} \quad e_1 \text{ used}, e_2 \text{ used}, \Delta \text{ avail}, \Lambda \downarrow V}{\vdash \Psi : \langle \alpha, e_1 \rangle, \Delta \Downarrow \langle \alpha^\perp, e_2 \rangle \parallel \Lambda : V} \\
[I_2] & \frac{\alpha \text{ atom} \quad e \text{ used}, \Delta \text{ avail}, \Lambda \downarrow V}{\vdash \alpha, \Psi : \Delta \Downarrow \langle \alpha^\perp, e \rangle \parallel \Lambda : V} \\
[D_1] & \frac{\phi \text{ not atom} \quad \vdash \Psi : \Delta \Downarrow \langle \phi, e \rangle \parallel e \text{ used}, \Lambda : V}{\vdash \Psi : \langle \phi, e \rangle, \Delta \Uparrow . \parallel \Lambda : V} \\
[D_2] & \frac{\phi \text{ not atom} \quad x \text{ new} \quad \vdash \Psi : \Delta \Downarrow \langle \phi, x \rangle \parallel x \text{ used}, \Lambda : V}{\vdash \phi, \Psi : \Delta \Uparrow . \parallel \Lambda : V}
\end{aligned}$$

(c) Identity and decide rules.

Figure 3: Focused constraint calculus for Linear Logic.

as a set of variable names such that the members have not yet occurred in the proof and are all distinct:

$$(\forall x_i \in X \mid x_i \text{ new}) \wedge (\forall x_i, x_j \in X \mid i \neq j \Rightarrow x_i \neq x_j)$$

Then the counter is updated as follows:

$$[\otimes] \frac{X \text{ new} \quad \vdash^{i+|X|,n} \Psi : \Delta_L^X \Downarrow \langle \phi_1, e \rangle \parallel e \text{ used}, \Lambda : V' \quad \vdash^{n,n'} \Psi : \Delta_R^X \Downarrow \langle \phi_2, e \rangle \parallel e \text{ used}, \Lambda : V'' \quad V' \sim_{\text{vars}(\Delta), X} V''}{\vdash^{i,n'} \Psi : \Delta \Downarrow \langle \phi_1 \otimes \phi_2, e \rangle \parallel \Lambda : V''}$$

Definition 2.3.11. Given a judgment

$$\vdash \Delta$$

the corresponding judgment in the calculus of Figure 3 is

$$\vdash . : . \Uparrow \hat{\Delta} \parallel X \text{ used} : V$$

with X new, and

$$\hat{\Delta} = \{\langle \phi_i, x_i \rangle \mid \phi_i \in \Delta, x_i \in X\}$$

for some arbitrary ordering of Δ and X .

Lemma 2.3.1. *For all sequents Δ and assignments V which cover $\text{vars}(\Delta) \cup X$:*

$$\Delta_L^X[V] \cap \Delta_R^X[V] = \emptyset$$

Proof. This is a simple consequence of the fact that if $\phi \in \Delta_L^X[V]$ there is a annotated formula $\langle \phi, e \rangle \in \Delta$ such that

$$e[V] = \top$$

Since e is defined as a conjunction on boolean variables, all the variables in it must evaluate to true. It is straightforward to see that if the variable added by the split in Δ_L^X is x_i , and the corresponding one in Δ_R^X is $\overline{x_i}$, then when x_i is true in the assignment V , x_i^\perp is false. Hence $\phi \notin \Delta_R^X[V]$. The same can be done to show that if $\phi \in \Delta_R^X[V]$ then $\phi \notin \Delta_L^X[V]$. ■

Lemma 2.3.2. *For all sequents Δ and assignments V which cover $\text{vars}(\Delta) \cup X$,*

$$\Delta[V] = \Delta_L^X[V] \cup \Delta_R^X[V]$$

Proof. The simpler side is $\Delta_L^X[V] \cup \Delta_R^X[V] \subseteq \Delta[V]$, since it holds by the definition of the split (Definition 2.3.9). For the other side, suppose there was a formula ϕ such that $\phi \in \Delta[V]$ and $\phi \notin \Delta_L^X[V] \cup \Delta_R^X[V]$. This means that for some variable x_i

$$\begin{aligned} \langle \phi, e \rangle \in \Delta &\Rightarrow e[V] = \top \\ \langle \phi, x_i \wedge e \rangle &\notin \Delta_L^X \Rightarrow x_i \wedge e[V] = \perp \\ \langle \phi, \overline{x_i} \wedge e \rangle &\notin \Delta_R^X \Rightarrow \overline{x_i} \wedge e[V] = \perp \end{aligned}$$

But either x_i or $\overline{x_i}$ must be true in a certain assignment, thus either $x_i \wedge e$ or $\overline{x_i} \wedge e$ must be true, contradicting the hypothesis. ■

Theorem 2.3.1 (Soundness). *For any judgment in our constraint calculus:*

- if $\vdash \Psi : \Delta \uparrow \Gamma \parallel \Lambda : V$, then $\vdash_A \Psi : \Delta \uparrow \Gamma$;
- if $\vdash \Psi : \Delta \downarrow \langle \phi, e \rangle \parallel \Lambda : V$, then $\vdash_A \Psi : \Delta \downarrow \phi$.

Where A is the triadic calculus of [2] (Figure 1).

Proof. We will proceed by mutual induction on the structure of the given derivation. The four base cases are proved essentially in the same way, with the exception of top (\top):

(I_1) : Given the rule [I_1]

$$[I_1] \frac{\alpha \text{ atom} \quad e_1 \text{ used}, e_2 \text{ used}, \Delta \text{ avail}, \Lambda \downarrow V}{\vdash \Psi : \langle \alpha, e_1 \rangle \downarrow \langle \alpha^\perp, e_2 \rangle, \Delta \parallel \Lambda : V}$$

looking at the constraints we get that

$$\begin{aligned}\Delta[V] &= \emptyset && \text{(Because of } \Delta \text{ avail)} \\ \langle \alpha, e_1 \rangle[V] &= \alpha && \text{(Because of } e_1 \text{ used)} \\ \langle \alpha^\perp, e_2 \rangle[V] &= \alpha^\perp && \text{(Because of } e_2 \text{ used)}\end{aligned}$$

so we can rewrite this as

$$[I_1] \frac{\alpha \text{ atom}}{\vdash_A \Psi : \alpha \Downarrow \alpha^\perp}$$

(I_2) : Given the rule [I_2]

$$[I_2] \frac{\alpha \text{ atom} \quad e \text{ used}, \Delta \text{ avail}, \Lambda \Downarrow V}{\vdash \Psi, \alpha : \Delta \Downarrow \langle \alpha^\perp, e \rangle \parallel \Lambda : V}$$

proceeding as above we get

$$\begin{aligned}\Delta[V] &= \emptyset \\ \langle \alpha^\perp, e \rangle[V] &= \alpha^\perp\end{aligned}$$

thus

$$[I_2] \frac{\alpha \text{ atom}}{\vdash_A \alpha, \Psi : . \Downarrow \alpha^\perp}$$

(1) : Given the rule [1]

$$[1] \frac{e \text{ used}, \Delta \text{ avail}, \Lambda \Downarrow V}{\vdash \Psi : \Delta \Downarrow \langle 1, e \rangle \parallel \Lambda : V}$$

proceeding as above we get

$$\begin{aligned}\Delta[V] &= \emptyset \\ \langle 1, e \rangle[V] &= 1\end{aligned}$$

thus

$$[1] \frac{}{\vdash_A \Psi : . \Downarrow 1}$$

(\top) : Given the rule [\top] we have that

$$\frac{}{\vdash \Psi : \Delta \Uparrow \langle \top, - \rangle, \Gamma \parallel - : -}$$

Thus we can choose whichever assignment V that covers $\text{vars}(\Delta, \Gamma)$, and obtain

$$\frac{}{\vdash_A \Psi : \Delta[V] \Uparrow \top, \Gamma[V]}$$

The induction then follows like this:

(\otimes) : Given the rule [\otimes], we apply the inductive hypothesis, and from the premises

$$\begin{aligned} \vdash \Psi : \Delta_L^X \Downarrow \langle \phi_1, e \rangle \parallel e \text{ used}, \Lambda : V' \\ \vdash \Psi : \Delta_R^X \Downarrow \langle \phi_2, e \rangle \parallel e \text{ used}, \Lambda : V'' \end{aligned}$$

we extract the triadic proofs

$$\begin{aligned} \vdash_A \Psi : \Delta_L^X[V'] \Downarrow \phi_1 \\ \vdash_A \Psi : \Delta_R^X[V''] \Downarrow \phi_2 \end{aligned}$$

Since $V' \sim_{\text{vars}(\Delta), X} V''$, this can be rewritten as

$$\begin{aligned} \vdash_A \Psi : \Delta_L^X[V''] \Downarrow \phi_1 \\ \vdash_A \Psi : \Delta_R^X[V''] \Downarrow \phi_2 \end{aligned}$$

Furthermore because of Lemma 2.3.1 we have that $\Delta_L^X[V'']$ and $\Delta_R^X[V'']$ are disjoint, so the contexts of the two branches are separated. Hence we can apply [\otimes]_A, and obtain

$$\vdash_A \Psi : \Delta_L^X[V''], \Delta_R^X[V''] \Downarrow \phi_1 \otimes \phi_2$$

which, because of Lemma 2.3.2, correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \Downarrow \langle \phi_1 \otimes \phi_2, e \rangle \parallel \Lambda : V''$$

should be mapped to.

($\&$) : Given the rule [$\&$], we apply the inductive hypothesis, and from the premises

$$\begin{aligned} \vdash \Psi : \Delta \Uparrow \langle \phi_1, e \rangle, \Gamma \parallel e \text{ used}, \Lambda : V' \\ \vdash \Psi : \Delta \Uparrow \langle \phi_2, e \rangle, \Gamma \parallel e \text{ used}, \Lambda : V'' \end{aligned}$$

we extract the triadic proofs

$$\begin{aligned} \vdash \Psi : \Delta[V'] \Uparrow \phi_1, \Gamma[V'] \\ \vdash \Psi : \Delta[V''] \Uparrow \phi_2, \Gamma[V''] \end{aligned}$$

Now, since $V' \sim_{\text{vars}(\Delta, \Gamma)} V''$, we can write

$$\begin{aligned} \vdash \Psi : \Delta[V''] \Uparrow \phi_1, \Gamma[V''] \\ \vdash \Psi : \Delta[V''] \Uparrow \phi_2, \Gamma[V''] \end{aligned}$$

this way we can apply $[\&]_A$ to obtain

$$\vdash_A \Psi : \Delta[V''] \uparrow \phi_1 \& \phi_2, \Gamma[V'']$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \uparrow \langle \phi_1 \& \phi_2, e \rangle, \Gamma \parallel \Lambda : V''$$

should be mapped to.

(\mathfrak{A}) : Given the rule $[\mathfrak{A}]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \Delta \uparrow \langle \phi_1, e \rangle, \langle \phi_2, e \rangle, \Gamma \parallel e \text{ used}, \Lambda : V$$

we extract the triadic proof

$$\vdash_A \Psi : \Delta[V] \uparrow \phi_1, \phi_2, \Gamma[V]$$

Then we apply $[\mathfrak{A}]_A$, and obtain

$$\vdash_A \Psi : \Delta[V] \uparrow \phi_1 \mathfrak{A} \phi_2, \Gamma[V]$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \uparrow \langle \phi_1 \mathfrak{A} \phi_2, e \rangle, \Gamma \parallel \Lambda : V$$

should be mapped to.

(\oplus_L) : Given the rule $[\oplus_L]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \Delta \Downarrow \langle \phi_1, e \rangle \parallel e \text{ used}, \Lambda : V$$

we extract the triadic proof

$$\vdash_A \Psi : \Delta[V] \Downarrow \phi_1$$

Then we apply $[\oplus_L]_A$, and obtain

$$\vdash_A \Psi : \Delta[V] \Downarrow \phi_1 \oplus \phi_2$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \Downarrow \langle \phi_1 \oplus \phi_2, e \rangle \parallel \Lambda : V$$

should be.

(\oplus_R) : Given the rule $[\oplus_R]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \Delta \Downarrow \langle \phi_2, e \rangle \parallel e \text{ used}, \Lambda : V$$

we extract the triadic proof

$$\vdash_A \Psi : \Delta[V] \Downarrow \phi_2$$

Then we apply $[\oplus_R]_A$, and obtain

$$\vdash_A \Psi : \Delta[V] \Downarrow \phi_1 \oplus \phi_2$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \Downarrow \langle \phi_1 \oplus \phi_2, e \rangle \parallel \Lambda : V$$

should be.

$(!)$: Given the rule $[\!]\!]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \Delta \Uparrow \langle \phi, e \rangle \parallel e \text{ used}, \Delta \text{ avail}, \Lambda : V$$

we extract the triadic proof

$$\vdash_A \Psi : . \Uparrow \phi$$

since $\Delta[V] = \emptyset$ under assignment V . We then apply $[\!]\!]_A$ and obtain

$$\vdash_A \Psi : . \Downarrow !\phi$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \Downarrow \langle !\phi, e \rangle \parallel \Lambda : V$$

should be mapped to

$(?)$: Given the rule $[\!]\!]$, we apply the inductive hypothesis, and from the premise

$$\vdash \phi, \Psi : \Delta \Uparrow \Gamma \parallel e \text{ used}, \Lambda : V$$

we extract the triadic proof

$$\vdash_A \phi, \Psi : \Delta[V] \Uparrow \Gamma[V]$$

Then we apply $[?]_A$ and obtain

$$\vdash_A \Psi : \Delta[V] \uparrow ?\phi, \Gamma[V]$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \uparrow \langle ?\phi, e \rangle, \Gamma \parallel \Lambda : V$$

should be mapped to.

(\perp) : Given the rule $[\perp]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \Delta \uparrow \Gamma \parallel e \text{ used}, \Lambda : V$$

we extract the triadic proof

$$\vdash_A \Psi : \Delta[V] \uparrow \Gamma[V]$$

Then we apply $[\perp]_A$ and obtain

$$\vdash_A \Psi : \Delta[V] \uparrow \perp, \Gamma[V]$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \uparrow \langle \perp, e \rangle, \Gamma \parallel \Lambda : V$$

should be mapped to.

(D_1) : Given the rule $[D_1]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \Delta \downarrow \langle \phi, e \rangle \parallel e \text{ used}, \Lambda : V$$

we extract the triadic proof

$$\vdash_A \Psi : \Delta[V] \downarrow \phi$$

Then we apply $[D_1]_A$ and obtain

$$\vdash_A \Psi : \phi, \Delta[V] \uparrow .$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \langle \phi, e \rangle, \Delta \uparrow . \parallel \Lambda : V$$

should be mapped to.

(D_2) : Given the rule $[D_2]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \Delta \Downarrow \langle \phi, x \rangle \parallel x \text{ used}, \Lambda : V$$

where x new, we extract the triadic proof

$$\vdash_A \Psi : \Delta[V] \Downarrow \phi$$

Then we apply $[D_2]_A$ and obtain

$$\vdash_A \phi, \Psi : \Delta[V] \Uparrow .$$

which correctly corresponds to what the conclusion

$$\vdash \phi, \Psi : \Delta \Uparrow . \parallel \Lambda : V$$

should be mapped to.

$(R\Uparrow)$: Given the rule $[R\Uparrow]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \langle \phi, e \rangle, \Delta \Uparrow \Gamma \parallel \Lambda : V$$

when we extract the triadic proof, two cases arise:

- ϕ disappears under assignment V , and we get

$$\vdash_A \Psi : \Delta[V] \Uparrow \Gamma[V]$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \Uparrow \langle \phi, e \rangle, \Gamma \parallel \Lambda : V$$

should be mapped to, since if ϕ disappears from the premise, it will also disappear from the conclusion.

- ϕ remains under assignment V , and we get

$$\vdash_A \Psi : \phi, \Delta[V] \Uparrow \Gamma[V]$$

Then we apply $[R\Uparrow]_A$ and obtain

$$\vdash_A \Psi : \Delta[V] \Uparrow \phi, \Gamma[V]$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \uparrow \langle \phi, e \rangle, \Gamma \parallel \Lambda : V$$

should be mapped to.

$(R\Downarrow)$: Given the rule $[R\Downarrow]$, we apply the inductive hypothesis, and from the premise

$$\vdash \Psi : \Delta \uparrow \langle \phi, e \rangle \parallel e \text{ used}, \Lambda : V$$

we extract the triadic proof

$$\vdash_A \Psi : \Delta[V] \uparrow \phi$$

Then we apply $[R\Downarrow]_A$, and obtain

$$\vdash_A \Psi : \Delta[V] \Downarrow \phi$$

which correctly corresponds to what the conclusion

$$\vdash \Psi : \Delta \Downarrow \langle \phi, e \rangle \parallel \Lambda : V$$

should be mapped to.

■

In Appendix A we show an example derivation using our calculus of Figure 3.

Chapter 3

Implementation

We now describe the main implementation details of our prover. During this section, to distinguish the variables of Definition 2.3.1 from Prolog's ones, we will always refer to the latter as Prolog variables. When explaining the code we will use some common names for Prolog variables:

- A is a set of unrestricted atoms, its purpose will be explained in Section 3.4.1;
- U is a set of unrestricted formulae; this more or less corresponds to Ψ from §2;
- F, F_1, \dots , are formulae, and Fs and D are a lists of them; these more or less correspond to ϕ, ϕ_1, Γ and Δ from §2;
- S is the queue of currently usable unrestricted formulae, its purpose will be explained in Section 3.4.2;
- L is a list of constraints; this more or less corresponds to Λ from §2.

3.1 Why Prolog

Prolog as a language and as an environment has been historically tied to automated theorem proving for its ability to express backtracking algorithms naturally. Most Prolog implementations also support CLP (constraint logic programming) through dedicated libraries. These allow to use constraints referencing some attributes of Prolog variables in the body of clauses; in our case we will use CLP(\mathcal{B}) [17], which provides tools to deal with boolean constraints. Its implementation is based on reduced and ordered Binary Decision Diagrams (BDDs). A boolean expression, in this context, is an expression made up of Prolog variables and the constants 1 and 0, respectively true and false. The allowed operators are the usual ones one would expect, in our case we will use exclusively conjunction, negation and equality, respectively:

```

X * Y.
~ X.
X ::= Y.

```

Usually constraints will be accumulated in a list; for this reason `CLP(\mathcal{B})` provides the functor `*(L)`, which is interpreted as the conjunction of the members of the list `L`. Finally the main predicate of the library is `sat/1`, which checks for the satisfiability of a constraint. Since we only deal with conjunctions we define the helper predicate `check/1`

```

check(L) :-
    sat(*(L)).

```

To better understand how the predicate works, we give some examples:

- if the constraints are not instantiated enough, they simply get reduced:

```

?- check([X * Y ::= 0]).
sat(1#X*Y).

```

Here `#` is exclusive or.

- if the constraints are unsatisfiable, the predicate fails:

```

?- check([X * Y ::= 1, X ::= 0]).
false.

```

- if the constraints are satisfiable the Prolog variables are unified to their assigned value:

```

?- check([X * Y ::= 0, X ::= 1]).
X = 1,
Y = 0.

```

Here one constraint (`X ::= 1`) even forces the other variable (`Y`) to be unified.

The automatic unification of Prolog variables to their value after constraint satisfaction is used to implicitly deal with the propagation of solutions: as soon as one is unified in one branch of the proof, its value will be propagated to any other constraint containing it.

We use the library with the flag `clpb_monotonic` set to `true`. This makes the algorithm many orders of magnitude faster with the condition that adding constraints cannot yield new solutions. As a side effect all Prolog variables must appear in a constraint wrapped by the functor `v/1`, so instead of writing

$X * Y ::= 1$

we have to write

$v(X) * v(Y) ::= 1$

This small explanation sums up the extent of the library we use in our prover.

3.2 Formula transformations

Before beginning the search a sequent passes through a number of transformations. These transformations both preprocess the sequent to a more convenient form, and also add information about the subformulae.

As a first transformation the sequent gets normalized into negated normal form (NNF) as defined in Definition 2.1.1. Normalization is a common technique used in all the provers we compare with. The implementation mirrors exactly the transformation from two-sided judgment to one-sided judgment of Section 2.1:

1. the left sequent is negated and appended to the right sequent;
2. a predicate, which encodes the DeMorgan rules, is mapped recursively over the new sequent.

From an implementation point of view the purpose of normalization is to reduce the number of available choices the prover has at a certain moment, although by doing so we sacrifice some of the structure of the formula.

Next, to each formula we assign its why-not height, a measure borrowed from another prover's implementation (Section 4.1).


Definition 3.2.1 (Why-not height). Why-not height is the maximum number of nested “why-not”s in a formula, or, for formulae as defined in Definition 2.1.1,

$$\text{wnh}(\phi) = \begin{cases} 0 & \text{if } \phi \in \{\perp, \top, 1, 0, \alpha, \alpha^\perp\} \\ \max(\text{wnh}(\phi_1), \text{wnh}(\phi_2)) & \text{if } \phi \in \{\phi_1 \otimes \phi_2, \phi_1 \wp \phi_2, \phi_1 \oplus \phi_2, \phi_1 \& \phi_2\} \\ \text{wnh}(\phi_1) & \text{if } \phi \in \{!\phi_1\} \\ 1 + \text{wnh}(\phi_1) & \text{if } \phi \in \{?\phi_1\} \end{cases}$$

The purpose of this attribute is to guide the prover to the reasonably simpler choice – the one with the least nested exponentials – at different times during proof search. This happens in three ways:

- When a rule generates two branches (\otimes , $\&$), the branch associated to the formula with the least why-not height is tried first. This will presumably make the prover choose the simpler branch of the two, making it so that we can fail early if the branch turns out to be false. An example of this is Section 3.4.
- In the case of plus (\oplus), the branch associated to the least why-not height is tried first; this makes it so we can continue if the branch turns out to be true, ignoring the other harder branch.
- During the D_2 rule, the exponentials are tried in order of ascending why-not height. This process is further explained in Section 3.4.2.

After this transformation, formulae are attribute trees with at each node the why-not height. For example the formula $*(a, ?(b))$ becomes $*(a-0, ?(b-0)-1)-1$, or



As a third and final transformation, each formula gets annotated as in Definition 2.3.2. This also returns the initial constraints, which set each formula in the sequent to “used”, thus stating that in the proof each and every formula must be used.

3.3 Helper predicates

For the prover we define a series of helper predicates ranging from ones that just implement “ ϕ asy” from Definition 2.2.1, to ones that aid the generation of constraints. In particular we now cover the implementation of the splitting function from Definition 2.3.9

```

1  %! split_ctx(+[AFs], -[AFs], -[AFs], -[Cns], -[Cns]) is det.
2  split_ctx(Afs, Pos, Neg, PCns, NCns) :-
3      maplist([ af(F, N, E)
4                , af(F, VarPos, Y)
5                , af(F, VarNeg, Z)
6                , v(Y) == v(X) * v(E)
7                , v(Z) == (~ v(X)) * v(E)

```

```

8         ]>>(
9             gensym(x, V),
10            atomic_list_concat([N, V], ' . ', VarPos),
11            atomic_list_concat([N, V], ' . ~ ', VarNeg)
12        ), Afs, Pos, Neg, PCns, NCns).

```

The code itself is nothing special, consisting of a simple map over the list of formulae generating the constraints. What is important is that the snippet above shows the clear distinction between variable names – represented by atoms (lines 10-11), and the value of a variable – represented by Prolog variables (lines 6-7). This separation of name and value is needed because – after checking the constraints – the solver unifies the Prolog variables to their values if it finds a valid assignment; the purpose of the atom is then to associate the variable value to its name if the final proof tree.

Finally, it must be noted that there is no predicate implementing the concept of two coinciding assignments (Definition 2.3.6): this – as mentioned before – is implicitly handled by Prolog’s unification mechanism.

3.4 Focusing

Focusing is implemented by two mutually defined predicates: `async/8` and `focus/8`, encoding respectively the asynchronous and the synchronous rules of Figure 3.

During the asynchronous phase there are two lists: `Fs` and `D`, representing Γ and Δ . The asynchronous rules always work on the head of `Fs`, breaking it down into its sub-formulae. This process can be seen for example in the clause for `with (&)`:

```

1  async(A, U, D, [F|Fs], S, M, L) :-
2      F = af(((F1-H1) & (F2-H2))-_), N, E), !,
3      ( H2 > H1
4      -> async(A, U, D, [af((F1-H1), N, E)|Fs], S, M, [v(E) == 1|L]),
5          async(A, U, D, [af((F2-H2), N, E)|Fs], S, M, [v(E) == 1|L])
6      ;  async(A, U, D, [af((F2-H2), N, E)|Fs], S, M, [v(E) == 1|L]),
7          async(A, U, D, [af((F1-H1), N, E)|Fs], S, M, [v(E) == 1|L])
8      ).

```

Here, we chose this particular rule to also show how the prover is guided using why-not height (line 3) as mentioned in Definition 3.2.1. Comparing this with the `[&]` rule from Figure 3, one can see that the implementation closely resembles the formal specification. The cut at line 2 is present in all the clauses for asynchronous connectives, and eliminates any choice point since these rules should be deterministic.

If a formula cannot be further be broken apart – i.e. it is either an atom, a negated atom, or it has a top-level synchronous connective – then it is put to the side in `D`. This process goes on as long as `Fs` is not empty.

When each formula is moved from **Fs** to **D**, the phase switches and the focusing process begins: a formula is chosen from either **D** or **U** by applying one of the decide rules. The decide rules will be discussed further ahead in Section 3.4.2. The chosen formula gets broken down until either an asynchronous connective or a negated atom is left. Unlike the asynchronous phase, the rules applied in this phase are non-deterministic and may be backtracked, so not cut is introduced. Other than this the implementation of the predicate `focus/8` is almost the same as `async/8`: the formula is broken down into its sub-formulae, and – if necessary – the why not heights are used to guide the proof search.

3.4.1 Identity rules

This process of alternating asynchronous and synchronous phases in normal focusing (i.e. without constraints) goes on until only a positive literal (in our case a negated atom) in **Fs** is left, and the corresponding negative literal (in our case just an atom) in either present in **U** or **D**. When this happens, the axioms – rules $[I_1]_A$ or $[I_2]_A$ – are applied to close the proof. In our case when the prover is focusing and it encounters a positive literal in **Fs**, it checks whether the corresponding negative literal exists in **D**. If this is true, then the helper predicate `set_to_zero/2` implementing “– avail” is used to generate the constraints for **D**, these are appended to the constraints gathered up to that point and then checked. A slightly different approach is taken for $[I_2]$. The rule would call for us to search a matching atom in Ψ , corresponding to **U**. Instead the prover keeps a set of only unrestricted atoms, called **A**, separated from the rest of the unrestricted formulae.

```

1 focus(A, U, D, F, _, _, L) :-
2   F = af(((~ T)-_), _, E),
3   is_atom(T),
4   member(T, A),
5   set_to_zero(D, Dz),
6   append([v(E) =:= 1|Dz], In, Cns),
7   check(Cns).
```

This small modification is based on the fact that once a negative literal is put to the side (be it in Δ or in Ψ), it stays there for the rest of the proof. So we chose to segregate the unrestricted atoms in a separate set, non polluting the one containing formulae which can be focused on.

3.4.2 Decide rules

For the decide rules, particularly for $[D_2]$, we use a modified version of another prover’s (Section 4.1) heuristic. The method consists of not using directly the set Ψ in the

$[D_2]$ rule, but instead keeping a queue of ordered unrestricted formulae which can be refilled only a certain number of times per-branch. This can be seen in the clause of `async/8` implementing the rule:

```

1  async(A, U, D, [], [H|T], M, L) :-
2    ( U \= []
3      -> gensym(x, X),
4          focus(A, U, D, af(H, X, E), T, M, [v(E) =:= 1|L])
5    ).
6  async(A, U, D, [], [_|T], M, L) :-
7    U \= [] -> async(A, U, D, [], T, M, L).
8  async(A, U, D, [], [], M, L) :-
9    ( U \= []
10     -> refill(U, M, S, M1),
11         early_stop(A, U, D, S, M1, L)
12    ).
```

Here the fifth argument is the queue and the sixth is the bound. Two cases arise:

- if the queue (*S*) is empty and the bound (*M*) is greater than zero then the set of unrestricted formulae *U* is taken and it is sorted based on why-not height. This can be seen in the predicate `refill/4`

```

1  refill(U, M, S, M1) :-
2    ( M \= 0, U \= []
3      -> sort(2, @=<, U, S),
4          M1 is M - 1
5    ).
```

The new sorted list of unrestricted formulae becomes the new *S* and *M* is decreased. Otherwise if *M* is 0 (line 2) the branch fails.

- if the queue is not empty, then the first formula in the queue is extracted (line 1-5) and added to the working set. Otherwise if the branch fails (line 6-7) the formula gets discarded and the next one in the queue is tried.

In particular, if the queue *S* is refilled, we do not directly call `async/8`, but instead call the predicate `early_stop/7` (line 10), defined as:

```

early_stop(A, U, D, [H|T], M, L) :-
  gensym(x, X),
  focus(A, U, D, af(H, X, E), T, M, [v(E) =:= 1|L]).
early_stop(A, U, D, [_|T], M, L) :-
  early_stop(A, U, D, T, M, L).
```

This is due to the fact that if the branch was not provable and we instead called directly `async/8` at line 10, the prover would try to refill the branch `M` times. What `early_stop/7` does is fail if the queue has just been refilled and it turns out the branch was not provable.

All of this is done to prevent loops introduced by the use of exponentials. For example, if no such measures were taken and the first formula in `U` was an asynchronous one and the prover is at a dead-end, the following loop would arise:

1. the prover contracts the asynchronous formula out of `U` using $[D_2]$;
2. the rule $[R\Downarrow]$ is applied;
3. the asynchronous formula is broken down until `Fs` is empty;
4. repeat from step 1.

This method forces the prover to try all the other unrestricted formulae before trying one a second time, and the bound completely eliminates the problem of loops, although obviously making some true sequents unprovable.

The rules $[D_1]$, $[I_1]$ are defined before the unrestricted counterparts, so that they are tried first.

3.5 Building the tree

In the listings above we omitted one parameter from the calls to `focus/8` and `async/8`, whose purpose is to accumulate the proof tree during the search. For example in the clause for `par` (\mathcal{P})

```
async(A, U, D, [F|Fs], S, M, L, node(par, A, U, D, [F|Fs], [T])) :-
  F = af(((F1 / F2)-_), N, E), !,
  Fs1 = [af(F1, N, E), af(F2, N, E)|Fs],
  async(A, U, D, Fs1, S, M, [v(E) := 1|L], T).
```

we can see clearly in the eight argument the structure of one node of the tree: a label, the context and an – optionally empty – list of sub-trees, here containing only the subtree `T`. A leaf is just a node with an empty list of sub-trees.

This term can be used in the end to reconstruct the actual proof tree by visiting it and – for each formula of each node – querying whether its variable is set to one, deleting it otherwise (this process is the same used in the proof for Theorem 2.3.1). A tree without the focusing infrastructure may be built by removing all the nodes regarding the phases (i.e. $R\Downarrow$ $R\Uparrow$ and decide rules) and by rebuilding the original sequent by appending `A`, `U`, `D` and `Fs` together as explained in [2]. A more sophisticated algorithm may even cancel out unwanted unrestricted formulae, that otherwise remain lingering in the sequent.

Chapter 4

Related work

Most bottom-up provers for classic linear logic use some combination of focusing and normalization to structure their proofs, with the notable exception of `llprover` [15] not using normalization. In this chapter we give a closer look to two such provers: the aforementioned `llprover` and `APLL` [18]. In particular we will talk about three main characteristics: normalization, how spitting is handled, and bounds for exponentials.

Obviously there are many other provers for linear logic, among all we cite:

- `symplic` [4], a top-down prover based on the inverse method proposed by K. Chaudhuri in [5];
- `linTAP` [12, 13], a tableau prover for MELL (the multiplicative and exponential fragment of linear logic).

Ultimately we chose to compare against the two provers above as they more or less implement the same algorithm as us.

4.1 APLL

`APLL` is the underlying prover of `click&collect` [3], an interactive tool for building linear logic proofs in sequent calculus. It provides 4 different searches – forward and backwards for classic and intuitionistic linear logic. We will focus on the backwards algorithm for classic linear logic.

The program is written in OCaml and implements a standard focused proof search on normalized formulae as seen in [11]. Sequent splitting when encountering a tensor is done by generating all the numbers up to $2^{|\Delta|}$ – where Δ is the sequent – and using the bit representation of those to create the two subsets. As we will see in Section 5.2 this implementation choice will result in a degradation of performance on formulae with a high number of multiplicatives.

Their particular usage of why-not height and implementation of the decide rule was a major influence for our own implementation and thus has been already discussed in §3. Still, for the sake of clarity, we repeat it here: the prover uses a fixed local bound; this bound does not represents the number of contraction per branch, but instead the number of times the unrestricted formulae can be copied in the queue per branch.

4.2 llprover

llprover is a prover written in Prolog by Naoyuki Tamura. Where APLL had different provers for intuitionistic and classical linear logic, this prover encodes all the rules as the same predicate `rule/6`, using the first argument as a selector for the system. Choosing full classical logic as the system uses all the rules, including the ones for intuitionistic linear logic. For this reason the prover does not use normalization; instead the option is given to transform the two-sided proof to a one-sided one.

Tensor splitting is implemented similarly to APLL by trying every possible partition

```

1 rule([ill,0], no, r(*), S, [S1, S2], [r(N),r(N1),r(N2)]) :-
2   match(S, ([X]-->[Y1,[A*B],Y2])),
3   merge(X1, X2, X),
4   merge(Y11, Y12, Y1),
5   merge(Y21, Y22, Y2),
6   match(S1, ([X1]-->[Y11,[A],Y21])),
7   match(S2, ([X2]-->[Y12,[B],Y22])),
8   length(Y1, N), length(Y11, N1), length(Y12, N2).
```

Here `merge/3` (lines 3, 4 and 5) – called with the only the third argument bound – generates all possible lists that when merged together return the original sequents.

Another particular characteristic of **llprover** is that it uses a local bound with iterative deepening: the prover will repeatedly try to prove the sequent, each time incrementing the bound up to the maximum specified. This guarantees finding the simplest proof, at the expense of the overall performance. This is unlike our prover or APLL, which instead directly use the specified bound.

Chapter 5

Testing

5.1 Infrastructure

The infrastructure for the project is divided in two parts:

- the Nix [1] “glue”,
- some python utilities to call and confront provers.

In nix everything must be packaged as something called a derivation: a declarative recipe specifying all the needed programs (as other derivations) and the steps necessary to build the source. These steps are executed in a environment containing only the specified programs and little else.

Derivations are written for all the programs we interact with:

- the provers, such as APLL or llprover;
- the formula generator adapted from APLL described in Section 5.1.3;
- the LLTP parser.

All these, together with one for the python environment, are then used to define a single development environment to run the Jupyter notebook with all the necessary dependencies. Furthermore, having used “flakes” (a nix experimental feature), all dependencies are locked to a certain commit thus ensuring better reproducibility.

The main logic for the benchmarking and testing is defined in the python module `testprover.py`. Here for simplicity we assume that the executable of a prover returns a code of 0 if it has found a proof, or any other number otherwise. This condition is already met by our prover and by llprover; for APLL we instead provide a wrapper.

Defining the call to our prover using this library is as simple as writing

```

1 pc = Registered()
2
3 @pc.register('sat-ll')
4 @call_prover(PrefixTree.SAT_LL_DICT)
5 def call_sat_prover(premises, conclusions):
6     return [ 'sat-ll'
7             , '-b', '3'
8             , f'{premises} /- {conclusions}'
9             ]

```

The function defining a call to a prover must take as inputs two arguments, the list of premises and the list of conclusions (using the format described in Section 5.1.1), and return the command calling the prover as a list. The innermost decorator (line 3) then accepts a dictionary, and returns a function which automatically calls the prover above with a test entry (as in Section 5.1.2), times it, and eventually terminates the process if it takes too much. Lastly, the outermost decorator (line 4) simply adds an entry 'sat-ll' associated with the prover call to the register `pc`. This register is just an association name to function, and it is what is actually passed to the benchmarking functions which will use the names in the output tables.

The library then provides two functions:

- **testall** accepts a single prover and a test suite; the function checks if the output of the prover corresponds to the expected value for the test;
- **benchmark** accepts a register and a test suite and returns the times and outcomes of each prover. This function is used to compare different provers.

All the outputs of the functions above are pandas **DataFrames**, which means that they can be easily queried, dumped to CSV, aggregated, or visualized using most plotting libraries. For example let `out` be the output of a call to **testall** with some prover and test suite:

```

out['outcome'].map(lambda x: x == Outcome.SUCCESS).all()
len(out[(out['outcome'] == Outcome.FAILURE)
        | (out['outcome'] == Outcome.TIMEOUT)])

```

test respectively whether all the tests succeeded and the number of failed tests.

5.1.1 Prefix format

Since for benchmarking we may interface with many different provers – each with its own syntax for expressions – the need arose for a common format which was easy

to parse and translate. For this purpose we define a prefix format for linear logic formulae inspired by the format used by [16] for implicational formulae:

formula	symbol
$\phi_A \otimes \phi_B$	*AB
$\phi_A \wp \phi_B$	 AB
$\phi_A \oplus \phi_B$	+AB
$\phi_A \& \phi_B$	&AB
$\phi_A \multimap \phi_B$	@AB
ϕ_A^\perp	^A
$?\phi_A$?A
$!\phi_A$!A

Each single character not representing an operator is considered as a variable name. Longer names can be specified by enclosing them in single quotes as in 'varname'. As an example we give the translation of DeMorgan for the tensor:

$$\text{prefix}((a \otimes b)^\perp \multimap a^\perp \wp b^\perp) = @^{\wedge}ab|^{\wedge}a^{\wedge}b$$

5.1.2 File formats

We use json as a standard format to store the tests, because of its vast adoption by most programming languages. A test suite is thus defined as a list of test cases; where a test case is just an object with these three mandatory fields:

id	is a number with the sole purpose of tracing back the test case from the output;
premises	is a list of premises as prefix formulae;
conclusions	is a list of conclusions as prefix formulae.

For example

```
{
  "id": 1,
  "premises": [ "^*ab" ],
  "conclusions": [ "/^a^b" ]
}
```

is the test case representing $(a \otimes b)^\perp \vdash a^\perp \wp b^\perp$. Other arbitrary fields may be present; for example we will use the following optional fields:

thm	tells whether this test case is a tautology or not, may be null;
*, +, ?, ...	is the number of times a specific connective appears in the test case;
notes	is a human readable string about the test case, e.g. its infix representation;
size	is an indicative number of the size of the formula;
atoms	is the upper bound on the number of atoms.

The fields **size** and **atoms** will be further explained in Section 5.1.3.

5.1.3 Formula generation

One of the sources of formulae we'll use in Section 5.2 is APLL's random formula generator. The version we'll use is a slight modification of it, where:

- the output is in the json format described in Section 5.1.2;
- one can choose whether to generate normalized formulae or not;
- one can choose which connectives appear in the generated formula.

A noteworthy detail is how the parameters **size** and **atoms** mentioned in Section 5.1.2 are defined, as these are directly related to how the formulae are generated:

- when one specifies a number of atoms, the generator initializes an array containing that number of atoms, their negations, and the constants \perp, \top, \dots . During the generation of the formula this array is randomly accessed, choosing an element when needed. This means that **atoms** represents an upper bound to the number of different atoms that may appear in the formula, not their exact number.
- when a formula is generated, at each step it is chosen whether to generate a unary or binary connective based on a threshold:
 - if a unary connective is chosen, the process continues with a size of

$$\text{size} - 1$$

- if a binary connective is chosen, the program chooses a random value between 0 and **size**, and it generates the two branches of the formula, with size respectively k and $\text{size} - k$.

This means that **size** is just an indicative value of the number of connectives in the formula.

5.2 Benchmarking

We'll mainly use three sources for formulae: llprover's tests; LLTP [14], especially the translations of Kleene's intuitionistic formulae; and randomly generated formulae made by the generator described in 5.1.3. Out of these, llprover's tests are composed primarily of simple linear logic tautologies, such as the DeMorgan rules; for this reason these tests are just used to quickly check for obvious bugs.

Randomly generated formulae will be used only for non-exponential test for the following reasons:

- datasets of linear logic theorems without exponentials are rare and most often the formulae in them do not have a significant size;
- when dealing with exponential formulae a prover may wrongfully deem a true formula false, just because it exhausted its bound. This problem is exacerbated when dealing with generated tests: since the expected output is not known we do not know if the prover terminated because of the bound, or if the test is actually a non-theorem.

Using random formulae we can see that our prover outperforms APLL (and llprover) when dealing with formulae rich in multiplicatives (Figure 4a and Table 2a).




Figure 4: Percentage of number of timeouts out of a hundred formulae

We can also see that in the multiplicative and additive case the difference begin to level (Figure 4b and Table 2b). The additive case is not that significant as the formulae remain manageable and no major differences can be seen. When reviewing

prover	timeouts	successes	success rate	avg. time (succ.)
APLL	34	66	0.66	1.441 s
llprover	50	50	0.50	3.006 s
sat-ll	7	93	0.93	1.874 s

(a) Data corresponding to Figure 4a

prover	timeouts	successes	success rate	avg. time (succ.)
APLL	24	76	0.76	1.213 s
llprover	42	58	0.58	0.807 s
sat-ll	7	93	0.93	1.289 s

(b) Data corresponding to Figure 4b

Table 2: Non exponential tests.

these outputs it is important to remember that randomly generated tests lack the structure of real formulae, and this may impact the quality of the measurement. All these test were done generating suites of 100 tests using a timeout of 60 s:

- Figure 4a’s tests were composed of formulae of **size** 100, and **atoms** 50, using just tensor (\otimes) and par (\wp) as connectives;
- Figure 4b’s tests were composed of formulae of **size** 500, and **atoms** 250, with all the connectives, except the exponentials (\otimes , \wp , $\&$ and \oplus)

prover	timeouts	failures	successes	success rate	avg. (succ.)	avg. (tot.)
APLL	0	17	71	≈ 0.80	0.035 s	0.326 s
llprover	20	6	62	≈ 0.70	0.981 s	2.179 s
sat-ll	5	15	68	≈ 0.77	0.443 s	0.496 s

(a) Outputs for KLE-cbv

prover	timeouts	failures	successes	success rate	avg. (succ.)	avg. (tot.)
APLL	0	16	72	≈ 0.80	0.037 s	0.055 s
llprover	20	6	62	≈ 0.70	1.709 s	3.253 s
sat-ll	4	18	66	≈ 0.75	0.130 s	0.185 s

(b) Outputs for KLE-cbn

Table 3: Exponential tests.

We now show the results of running the provers on two datasets: KLE-cbn and KLE-cbv, respectively the call-by-name and call-by-value translations of Kleene’s theorems. These translations introduce a high number of exponentials; this causes

– as said before – some failures not due to bugs, but because of the exhausted number of contractions. The benchmarks are done using a timeout of 60 s and a bound of 3. It can be seen in Table 3a and 3b that the differences of the previous tests are completely leveled, and instead our prover performs slightly worse than APLL.

Overall, since *llprover* uses incremental search, its times are often skewed towards the slower side. Similarly our prover is consistently slightly slower than APLL, this difference is negligible and can be probably ascribed to differences languages' run-times.

Chapter 6

Conclusion

The purpose of the original paper [9] was to show a general way to handle splitting which could be adapted to a series of different logics. In §5 we show that, after modifying the original calculus using techniques from proof theory (i.e. focusing [2] and normalization), this usage of constraints is an effective way of handling multiplicatives during linear logic proof search. Furthermore in §3, our Prolog implementation demonstrates how this method can be fairly easily integrated into a standard focusing prover without needing excessive modification, in part thanks to Prolog's elegant interface for CLP.

There are many ways in which our work may be expanded in the future, just to cite a few:

- our prover as of right now handles exponentials rather poorly (see Section 5.2), research could be made to devise alternative heuristics or a more refined use of constraints;
- a constraint calculus for intuitionistic linear logic (ILL) may be defined and a prover similar to ours implemented; the same could be said about the logic of bunched implications (BI) using the calculus described in [9];
- more exotic uses of constraints may be explored, perhaps using disjunction on the axioms such that

$$\vdash \langle a, e_1 \rangle, \langle a, e_2 \rangle, \langle a^\perp, e_3 \rangle$$

generates the constraints

$$((e_1 \text{ used} \wedge e_2 \text{ avail}) \vee (e_1 \text{ avail} \wedge e_2 \text{ used})) \wedge e_3 \text{ used}$$

Appendix A

Example derivation

The proof to the judgment

$$(a \otimes b)^\perp \vdash a^\perp \wp b^\perp$$

which is normalized to

$$\vdash (a \otimes b), a^\perp \wp b^\perp$$

corresponds in our calculus to

$$\begin{array}{c} \frac{[\otimes] \quad \frac{\nabla'}{\vdash . : \langle a^\perp, x_2 \rangle, \langle b^\perp, x_2 \rangle \Downarrow \langle a \otimes b, x_1 \rangle \parallel x_1 \text{ used}, x_2 \text{ used} : V} \quad \nabla''}{[D_1] \quad \vdash . : \langle a^\perp, x_2 \rangle, \langle b^\perp, x_2 \rangle, \langle a \otimes b, x_1 \rangle \Uparrow . \parallel x_1 \text{ used}, x_2 \text{ used} : V} \\ [R\Uparrow] \quad \vdash . : \langle b^\perp, x_2 \rangle, \langle a \otimes b, x_1 \rangle \Uparrow \langle a^\perp, x_2 \rangle \parallel x_1 \text{ used}, x_2 \text{ used} : V \\ [R\Uparrow] \quad \vdash . : \langle a \otimes b, x_1 \rangle \Uparrow \langle b^\perp, x_2 \rangle, \langle a^\perp, x_2 \rangle \parallel x_1 \text{ used}, x_2 \text{ used} : V \\ [\wp] \quad \vdash . : \langle a \otimes b, x_1 \rangle \Uparrow \langle a^\perp \wp b^\perp, x_2 \rangle \parallel x_1 \text{ used}, x_2 \text{ used} : V \\ [R\Uparrow] \quad \vdash . : . \Uparrow \langle a \otimes b, x_1 \rangle, \langle a^\perp \wp b^\perp, x_2 \rangle \parallel x_1 \text{ used}, x_2 \text{ used} : V \end{array}$$

with

$$\begin{array}{c} \nabla' = \frac{[I_1] \quad \frac{x_1 \text{ used}, x_2 \text{ used}, x_2 x_3 \text{ used}, x_2 x_4 \text{ avail} \Downarrow V}{\vdash . : \langle a, x_1 \rangle, \langle \cancel{b^\perp, x_2 x_4} \rangle \Downarrow \langle a^\perp, x_2 x_3 \rangle \parallel x_1 \text{ used}, x_2 \text{ used} : V}}{[D_1] \quad \vdash . : \langle a, x_1 \rangle, \langle a^\perp, x_2 x_3 \rangle, \langle \cancel{b^\perp, x_2 x_4} \rangle \Uparrow . \parallel x_1 \text{ used}, x_2 \text{ used} : V} \\ [R\Uparrow] \quad \vdash . : \langle a^\perp, x_2 x_3 \rangle, \langle \cancel{b^\perp, x_2 x_4} \rangle \Uparrow \langle a, x_1 \rangle \parallel x_1 \text{ used}, x_2 \text{ used} : V \\ [R\Downarrow] \quad \vdash . : \langle a^\perp, x_2 x_3 \rangle, \langle \cancel{b^\perp, x_2 x_4} \rangle \Downarrow \langle a, x_1 \rangle \parallel x_1 \text{ used}, x_2 \text{ used} : V \\ \nabla'' = \frac{[I_1] \quad \frac{x_1 \text{ used}, x_2 \text{ used}, x_3 \text{ used}, x_4 \text{ avail}, x_2 \bar{x}_4 \text{ used}, x_2 \bar{x}_3 \text{ avail} \Downarrow V}{\vdash . : \langle b, x_1 \rangle, \langle \cancel{a^\perp, x_2 \bar{x}_3} \rangle \Downarrow \langle b^\perp, x_2 \bar{x}_4 \rangle \parallel x_1 \text{ used}, x_2 \text{ used}, x_3 \text{ used}, x_4 \text{ avail} : V}}{[D_1] \quad \vdash . : \langle b, x_1 \rangle, \langle \cancel{a^\perp, x_2 \bar{x}_3} \rangle, \langle b^\perp, x_2 \bar{x}_4 \rangle \Uparrow . \parallel x_1 \text{ used}, x_2 \text{ used}, x_3 \text{ used}, x_4 \text{ avail} : V} \\ [R\Uparrow] \quad \vdash . : \langle \cancel{a^\perp, x_2 \bar{x}_3} \rangle, \langle b^\perp, x_2 \bar{x}_4 \rangle \Uparrow \langle b, x_1 \rangle \parallel x_1 \text{ used}, x_2 \text{ used}, x_3 \text{ used}, x_4 \text{ avail} : V \\ [R\Downarrow] \quad \vdash . : \langle \cancel{a^\perp, x_2 \bar{x}_3} \rangle, \langle b^\perp, x_2 \bar{x}_4 \rangle \Downarrow \langle b, x_1 \rangle \parallel x_1 \text{ used}, x_2 \text{ used}, x_3 \text{ used}, x_4 \text{ avail} : V \end{array}$$

and

$$V = \{x_1 \mapsto \top, x_2 \mapsto \top, x_3 \mapsto \top, x_4 \mapsto \perp\}$$

For reference we show the classic proof (non focused, without constraints) for the same judgment:

$$\begin{array}{c} [A] \frac{}{\vdash a, a^\perp} \quad [A] \frac{}{\vdash b, b^\perp} \\ [\otimes] \frac{}{\vdash a^\perp, b^\perp, (a \otimes b)} \\ [\wp] \frac{}{\vdash (a \otimes b), (a^\perp \wp b^\perp)} \end{array}$$

Bibliography

- [1] Nix & NixOS: Declarative builds and deployments. <https://nixos.org>. Accessed: 2024-06-27.
- [2] Jean-Marc Andreoli. Logic programming with focusing proofs in linear logic. *J. Log. Comput.*, 2(3):297–347, 1992.
- [3] Etienne Callies and Olivier Laurent. Click and coLLecT an interactive linear logic prover. In *5th International Workshop on Trends in Linear Logic and Applications (TLLA 2021)*, 2021.
- [4] Kaustuv Chaudhuri. A propositional linear inverse method theorem prover. <https://github.com/chaudhuri/sympli>. Accessed: 2024-03-20.
- [5] Kaustuv Chaudhuri and Frank Pfenning. Focusing the inverse method for linear logic. In C.-H. Luke Ong, editor, *Computer Science Logic, 19th International Workshop, CSL 2005, 14th Annual Conference of the EACSL, Oxford, UK, August 22-25, 2005, Proceedings*, volume 3634 of *Lecture Notes in Computer Science*, pages 200–215. Springer, 2005.
- [6] Gerhard Gentzen. Untersuchungen über das logische schließen. i. *Mathematische Zeitschrift*, 39(1):176–210, Dec 1935.
- [7] Gerhard Gentzen. Untersuchungen über das logische schließen. ii. *Mathematische Zeitschrift*, 39(1):405–431, Dec 1935.
- [8] Jean-Yves Girard. Linear logic. *Theor. Comput. Sci.*, 50:1–102, 1987.
- [9] James Harland and David J. Pym. Resource-distribution via boolean constraints. *ACM Trans. Comput. Log.*, 4(1):56–90, 2003.
- [10] Joshua S. Hodas and Dale Miller. Logic programming in a fragment of intuitionistic linear logic. *Inf. Comput.*, 110(2):327–365, 1994.
- [11] Chuck C. Liang and Dale Miller. Focusing and polarization in linear, intuitionistic, and classical logics. *Theor. Comput. Sci.*, 410(46):4747–4768, 2009.

- [12] Heiko Mantel and Jens Otten. lintap: A tableau prover for linear logic. <https://www.leancop.de/lintap>. Accessed: 2024-03-20.
- [13] Heiko Mantel and Jens Otten. lintap: A tableau prover for linear logic. In Neil V. Murray, editor, *Automated Reasoning with Analytic Tableaux and Related Methods, International Conference, TABLEUX '99, Saratoga Springs, NY, USA, June 7-11, 1999, Proceedings*, volume 1617 of *Lecture Notes in Computer Science*, pages 217–231. Springer, 1999.
- [14] Carlos Olarte, Giselle Reis, Elaine Pimentel, Valeria de Paiva, and Olivier Laurent. Linear logic theorem proving. <https://github.com/meta-logic/lltp>. Accessed: 2024-06-27.
- [15] Naoyuki Tamura. A linear logic prover (llprover). <https://cspsat.gitlab.io/llprover>. Accessed: 2024-03-20.
- [16] Paul Tarau and Valeria de Paiva. Deriving theorems in implicative linear logic, declaratively. In Francesco Ricca, Alessandra Russo, Sergio Greco, Nicola Leone, Alexander Artikis, Gerhard Friedrich, Paul Fodor, Angelika Kimmig, Francesca A. Lisi, Marco Maratea, Alessandra Mileo, and Fabrizio Riguzzi, editors, *Proceedings 36th International Conference on Logic Programming (Technical Communications), ICLP Technical Communications 2020, (Technical Communications) UNICAL, Rende (CS), Italy, 18-24th September 2020*, volume 325 of *EPTCS*, pages 110–123, 2020.
- [17] Markus Triska. Boolean constraints in SWI-Prolog: A comprehensive system description. *Science of Computer Programming*, 164:98 – 115, 2018. Special issue of selected papers from FLOPS 2016.
- [18] Jui-Hsuan Wu. A linear logic prover implemented in ocaml. https://github.com/wujuihsuan2016/LL_prover. Accessed: 2024-03-20.